

# Bidirectional LSTM with pre-trained Twitter Word Embeddings

Cliche used an ensemble of bidirectional LSTMs along with CNNs to produce state of the art results in Twitter sentiment analysis. He trains initial word embeddings on a large, unlabeled corpus of Twitter data using a neural language model. We will instead be using Stanford's pre-trained Glove word embeddings that were specifically trained on Twitter data. Since our training data is not very large, we anticipate that using these pre-trained word embeddings will result in an increase in performance.

```
In [2]: import pandas as pd
import re
import nltk
import string
import os
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from nltk.tokenize import word_tokenize, sent_tokenize
from nltk.stem.wordnet import WordNetLemmatizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.preprocessing.text import Tokenizer
from sklearn.model_selection import train_test_split
import glob, os
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Embedding, Dense
import numpy as np
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, f1_score
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:516: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:517: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
```

```

in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:518: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:525: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
np_resource = np.dtype(["resource", np.ubyte, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:541: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:542: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:543: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:544: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.

```

```

_np_quint16 = np.dtype([("quint16", np.uint16, 1)])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:545: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
_np_qint32 = np.dtype([("qint32", np.int32, 1)])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:550: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
np_resource = np.dtype([("resource", np.ubyte, 1)])
Using TensorFlow backend.

```

```
In [3]: os.chdir("data/")
```

Helper methods for reading tweets and cleaning them.

```

In [43]: def read_tsv(file_path):
          df = pd.read_table(file_path)
          return df

import string
import re

# code inspired from https://www.kaggle.com/rahulvv/bidirectional-lstm
-glove200d

def remove_urls(text):
    url = re.compile(r'https?://\S+|www\.\S+')
    return url.sub(r'', text)

def remove_html(text):
    html=re.compile(r'<.*?>')
    return html.sub(r'', text)

def split_text(text):
    text = text.split()
    return text

def lower(text):
    text = [word.lower() for word in text]
    return str(text)

def remove_punct(text):
    text = ''.join([char for char in text if char not in string.punctu

```

```

ation])
    text = re.sub('[0-9]+', '', str(text))
    return text

def remove_stopwords(text):
    pattern = re.compile(r'\b( '+r'|'.join(stopwords.words('english'))
+ r')\b\s*')
    text = pattern.sub(' ', text)
    return text

lemmatizer = WordNetLemmatizer()
def lemmatize_words(text):
    text = lemmatizer.lemmatize(text)
    return text

def clean_tweet(text):
    t0 = remove_urls(text)
    t1 = remove_html(t0)
    t2 = split_text(t1)
    t3 = lower(t2)
    t4 = remove_punct(t3)
    t5 = remove_stopwords(t4)
    t6 = lemmatize_words(t5)
    return t6

```

```

In [44]: tweet_df = pd.DataFrame(columns=['tweet', 'sentiment', 'NA'])
df_test = pd.DataFrame(columns=['tweet', 'sentiment', 'NA'])

for file in glob.glob("*.tsv"):
    if 'final_test' in file:
        df_test_cur = read_tsv(file)
        df_test = pd.concat([df_test, df_test_cur])
    else:
        df_train_cur = read_tsv(file)
        tweet_df = pd.concat([tweet_df, df_train_cur])

```

```
In [45]: print(tweet_df[['tweet', 'sentiment']])
```

	tweet	sentiment
0	05 Beat it - Michael Jackson - Thriller (25th ...	neutral
1	Jay Z joins Instagram with nostalgic tribute t...	positive
2	Michael Jackson: Bad 25th Anniversary Edition ...	neutral
3	I liked a @YouTube video http://t.co/AaR3pjp2P...	positive
4	18th anniv of Princess Diana's death. I still ...	positive
...	...	...
1137	Maybe it was - his - fantasy ?	positive
1138	It was ok , but they always just seem so nervo...	negative
1139	It is streamable from YepRoc -- matter of fact...	positive
1140	comment telling me who you are , or how you fo...	positive
1141	im on myspace ... ill try and find you and add...	neutral

[53368 rows x 2 columns]

```
In [46]: print(df_test[['tweet', 'sentiment']])
```

	tweet	sentiment
0	#ArianaGrande Ari By Ariana Grande 80% Full ht...	neutral
1	Ariana Grande KIIS FM Yours Truly CD listening...	positive
2	Ariana Grande White House Easter Egg Roll in W...	positive
3	#CD #Musics Ariana Grande Sweet Like Candy 3.4...	positive
4	SIDE TO SIDE 🙄 @arianagrande #sidetoside #aria...	neutral
...	...	...
11901	@dansen17 update: Zac Efron kissing a puppy ht...	positive
11902	#zac efron sex pic skins michelle sex https://...	neutral
11903	First Look at Neighbors 2 with Zac Efron Shirt...	neutral
11904	zac efron poses nude #lovely libra porn https:...	neutral
11905	#Fashion #Style The Paperboy (NEW Blu-ray Disc...	neutral

[11906 rows x 2 columns]

Reading Glove word embeddings into a dictionary.

```
In [47]: #preparing train lables
tweet_df.loc[tweet_df.sentiment == "positive", "sentiment"] = 2
tweet_df.loc[tweet_df.sentiment == "neutral", "sentiment"] = 1
tweet_df.loc[tweet_df.sentiment == "negative", "sentiment"] = 0

labels = tweet_df["sentiment"].tolist()
labels = [ int(x) for x in labels ]

#preparing test labels
df_test.loc[df_test.sentiment == "positive", "sentiment"] = 2
df_test.loc[df_test.sentiment == "neutral", "sentiment"] = 1
df_test.loc[df_test.sentiment == "negative", "sentiment"] = 0

labels_test = df_test["sentiment"].tolist()
labels_test = [ int(x) for x in labels_test ]
```

Converting tweets and labels into lists.

```
In [48]: train_tweets = tweet_df.tweet.values
y_train_orig = tweet_df.sentiment.values
test_tweets = df_test.tweet.values
```

```
In [49]: from keras.utils import to_categorical

train_labels = to_categorical(y_train_orig)

clean_training_tweets = []
for i in range(len(train_tweets)):
    data = clean_tweet(train_tweets[i])
    clean_training_tweets.append(data)

clean_testing_tweets = []
for i in range(len(test_tweets)):
    data = clean_tweet(test_tweets[i])
    clean_testing_tweets.append(data)
```

Checking the tweets after cleaning them.

```
In [50]: print(clean_training_tweets[:10])
print(clean_testing_tweets[:10])
```

```
[' beat michael jackson thriller th anniversary edition hd', 'jay
z joins instagram nostalgic tribute michael jackson jay z apparent
ly joined instagram saturday ', 'michael jackson bad th anniversar
y edition picture vinyl unique picture disc vinyl includes origina
l ', ' liked youtube video one direction singing man mirror mich
ael jackson atlanta ga june ', 'th anniv princess dianas death st
ill want believe living private island away public michael j
ackson', 'oridaganjazz st time heard michael jackson sing honolu
lu hawaii restaurant radio abc loved ', 'michael jackson ap
peared saturday th place top miamis trends trndnl', ' old en
ough remember michael jackson attending grammys brooke shields w
ebster sat lap show', 'etbrowser u enjoy nd rate michael jackso
n bit honest ques like cant feel face song god obvious want mj
', ' weeknd closest thing may get michael jackson long timeesp
ecially since damn near mimics everything']
['arianagrande ari ariana grande full singer actress', 'ariana gra
nde kiis fm truly cd listening party burbank arianagrande', 'arian
a grande white house easter egg roll washington arianagrande', 'cd
musics ariana grande sweet like candy oz ml sealed box authentic
new', 'side side 🥺 arianagrande sidetoside arianagrande musically
comunidadgay lgbt🌈 lotb...', 'hairspray live previews macys thanksg
iving day parade arianagrande televisionnbc', 'lindsaylohan 'feelin
g thankful' blasting arianagrande wearing 'toomuch...', ' hate lo
ve songs dammit arianagrande', 'ariana grande [right ft big sean]
アリアナ arianagrande', ' one would prefer listen whole day 🥰👉
could never choose arianagrande intoyou sidetoside songs poll']
```

```
In [20]: print('Loading word vectors...')
word2vec = {}
with open(os.path.join('../glove/glove.twitter.27B.200d.txt'), encodin
g = "utf-8") as f:
    for line in f:
        values = line.split()
        word = values[0]
        vec = np.asarray(values[1:], dtype='float32')
        word2vec[word] = vec
print('Found %s word vectors.' % len(word2vec))
```

```
Loading word vectors...
Found 1193514 word vectors.
```

```
In [52]: # converting tweets to integer sequences
tokenizer = Tokenizer(num_words= 20000, oov_token= 'OOV')
tokenizer.fit_on_texts(clean_training_tweets)
train_tweet_sequences = tokenizer.texts_to_sequences(clean_training_tweets)
word_index_train = tokenizer.word_index
print('Found %s unique words in train tweets.' % len(word_index_train))
X_train = pad_sequences(sequences=train_tweet_sequences, maxlen=32, padding= 'post', truncating='post')

test_tweet_sequences = tokenizer.texts_to_sequences(clean_testing_tweets)
X_test = pad_sequences(sequences= test_tweet_sequences, maxlen=32, padding='post', truncating='post')
```

Found 67101 unique words in train tweets.

```
In [53]: print('Shape of X train tensor: ', X_train.shape)
print('Shape of X test: ', X_test.shape)
```

Shape of X train tensor: (53368, 32)  
Shape of X test: (11906, 32)

```
In [54]: num_words = min(20000, len(word_index_train)+1)
embedding_matrix = np.zeros((num_words, 200))

embeddings = []
for word, i in word_index_train.items():
    if i<20000:
        embeddings = word2vec.get(word)
        if embeddings is not None:
            embedding_matrix[i] = embeddings
```

```
In [17]: model = tf.keras.Sequential()
model.add(tf.keras.layers.Embedding(input_dim=num_words,output_dim = 200, weights=[embedding_matrix], input_length=32,trainable=False))
model.add(tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(100, return_sequences=True)))
model.add(tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(3, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer=tf.keras.optimizers.Adam(lr=0.01), metrics=[ 'accuracy'])
```



WARNING:tensorflow:From /Users/shivaomrani/opt/anaconda3/envs/neural\_networks/lib/python3.7/site-packages/tensorflow/python/keras/initializers.py:119: calling RandomUniform.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /Users/shivaomrani/opt/anaconda3/envs/neural\_networks/lib/python3.7/site-packages/tensorflow/python/ops/init\_ops.py:1251: calling VarianceScaling.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /Users/shivaomrani/opt/anaconda3/envs/neural\_networks/lib/python3.7/site-packages/tensorflow/python/ops/init\_ops.py:97: calling GlorotUniform.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /Users/shivaomrani/opt/anaconda3/envs/neural\_networks/lib/python3.7/site-packages/tensorflow/python/ops/init\_ops.py:97: calling Orthogonal.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /Users/shivaomrani/opt/anaconda3/envs/neural\_networks/lib/python3.7/site-packages/tensorflow/python/ops/init\_ops.py:97: calling Zeros.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

In [18]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 32, 200)	4000000
bidirectional (Bidirectional)	(None, 32, 200)	240800
bidirectional_1 (Bidirectional)	(None, 32, 64)	59648
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 3)	6147
Total params: 4,306,595		
Trainable params: 306,595		
Non-trainable params: 4,000,000		

In [19]: `history=model.fit(X_train, train_labels, batch_size=128, epochs=15)`

WARNING:tensorflow:From /Users/shivaomrani/opt/anaconda3/envs/neural\_networks/lib/python3.7/site-packages/tensorflow/python/ops/math\_grad.py:1250: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Epoch 1/15

53368/53368 [=====] - 103s 2ms/sample - loss: 0.7818 - acc: 0.6416

Epoch 2/15

53368/53368 [=====] - 108s 2ms/sample - loss: 0.6944 - acc: 0.6876

Epoch 3/15

53368/53368 [=====] - 100s 2ms/sample - loss: 0.6321 - acc: 0.7211

Epoch 4/15

53368/53368 [=====] - 99s 2ms/sample - loss: 0.5661 - acc: 0.7535

Epoch 5/15

53368/53368 [=====] - 95s 2ms/sample - loss: 0.4978 - acc: 0.7854

Epoch 6/15

53368/53368 [=====] - 92s 2ms/sample - loss: 0.4374 - acc: 0.8170

Epoch 7/15

```

53368/53368 [=====] - 91s 2ms/sample - loss
: 0.3940 - acc: 0.8358
Epoch 8/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.3537 - acc: 0.8544
Epoch 9/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.3260 - acc: 0.8683
Epoch 10/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.3010 - acc: 0.8809
Epoch 11/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.2785 - acc: 0.8896
Epoch 12/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.2758 - acc: 0.8918
Epoch 13/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.2516 - acc: 0.9013
Epoch 14/15
53368/53368 [=====] - 92s 2ms/sample - loss
: 0.2503 - acc: 0.9038
Epoch 15/15
53368/53368 [=====] - 94s 2ms/sample - loss
: 0.2383 - acc: 0.9092

```

```
In [20]: pred_p = model.predict(X_test)
```

```
In [21]: pred = (np.round(pred_p)).astype(int)
final_pred = []
for sample in pred:
    pred_label = sample.argmax()
    final_pred.append(pred_label)
```

```
In [22]: y_binary = to_categorical(labels_test)
model.evaluate(x = X_test, y = y_binary )
```

```

11906/11906 [=====] - 10s 830us/sample - loss: 1.6071 - acc: 0.5879

```

```
Out[22]: [1.6070550479709491, 0.58793885]
```

```
In [23]: from sklearn.metrics import classification_report
print(classification_report(labels_test, final_pred))
```

	precision	recall	f1-score	support
0	0.58	0.61	0.59	3811
1	0.62	0.58	0.60	5743
2	0.51	0.56	0.53	2352
accuracy			0.58	11906
macro avg	0.57	0.58	0.58	11906
weighted avg	0.59	0.58	0.58	11906

```
In [24]: # Calling `save('my_model')` creates a SavedModel folder `my_model`.
model.save("bidirectional-lstm")
```

```
In [56]: # It can be used to reconstruct the model identically.
reconstructed_model = keras.models.load_model("bidirectional-lstm")
y_binary = to_categorical(labels_test)
reconstructed_model.evaluate(x = X_test, y =y_binary)
```

```
11906/11906 [=====] - 13s 1ms/sample - loss
: 1.6071 - acc: 0.5879
```

```
Out[56]: [1.6070550479709491, 0.58793885]
```

```
In [60]: from sklearn.metrics import classification_report
pred_p = reconstructed_model.predict(X_test)

pred = (np.round(pred_p)).astype(int)
final_pred = []
for sample in pred:
    pred_label = sample.argmax()
    final_pred.append(pred_label)

print(classification_report(labels_test, final_pred))
```

	precision	recall	f1-score	support
0	0.58	0.61	0.59	3811
1	0.62	0.58	0.60	5743
2	0.51	0.56	0.53	2352
accuracy			0.58	11906
macro avg	0.57	0.58	0.58	11906
weighted avg	0.59	0.58	0.58	11906

In [ ]: