

Convolutinal Neural Network for Twitter Sentiment Analysis

As Kim suggested in his 2014 paper, Convolutional neural networks for sentence classification, a simple CNN with a single convolutional layer can produce excellent results for sentence classification. Therefore, we decided to test their effectiveness for our task. In this notebook, we will train our own word embeddings on the training dataset. The embedding layer is then followed by a convolution layer.

```
In [1]: import pandas as pd
import re
from keras.preprocessing.text import Tokenizer
import torch
import torch.nn as nn
import torch.nn.functional as F
from sklearn.model_selection import train_test_split
import glob, os
```

Using TensorFlow backend.

```
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:516: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:517: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:518: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)])
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
```

```

site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
    _np_qint32 = np.dtype [("qint32", np.int32, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorflow/python/framework/dtypes.py:525: FutureWarni
ng: Passing (type, 1) or 'ltype' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
    np_resource = np.dtype [("resource", np.ubyte, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:541: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
    _np_qint8 = np.dtype [("qint8", np.int8, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:542: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
    _np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:543: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
    _np_qint16 = np.dtype [("qint16", np.int16, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:544: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
    _np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:545: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
    _np_qint32 = np.dtype [("qint32", np.int32, 1)]
/Users/shivaomrani/opt/anaconda3/envs/neural_networks/lib/python3.7/
site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:550: Futu
reWarning: Passing (type, 1) or 'ltype' as a synonym of type is depr
ecated; in a future version of numpy, it will be understood as (type
, (1,)) / '(1,)type'.
    np_resource = np.dtype [("resource", np.ubyte, 1)]

```

```
In [2]: os.chdir("data/")
```

Helper functions for reading training and test datasets into dataframes and for pre-processing tweets.

```
In [3]: from nltk.stem.wordnet import WordNetLemmatizer

# inspired from https://www.kaggle.com/rahulvv/bidirectional-lstm-glove200d

def read_tsv(file_path):
    df = pd.read_table(file_path)
    return df

def remove_urls(text):
    url = re.compile(r'https?://\S+|www\.\S+')
    return url.sub('', text)

def remove_html(text):
    html = re.compile(r'<.*?>')
    return html.sub('', text)

def split_text(text):
    text = text.split()
    return text

def lower(text):
    text = [word.lower() for word in text]
    return str(text)

def remove_punct(text):
    text = ''.join([char for char in text if char not in string.punctuation])
    text = re.sub('[0-9]+', '', str(text))
    return text

def remove_stopwords(text):
    pattern = re.compile(r'\b(' + r'|'.join(stopwords.words('english')) + r')\b\s*')
    text = pattern.sub(' ', text)
    return text

lemmatizer = WordNetLemmatizer()
def lemmatize_words(text):
    text = lemmatizer.lemmatize(text)
    return text
```

```
def clean_tweet(text):
    t0 = remove_urls(text)
    t1 = remove_html(t0)
    t2 = split_text(t1)
    t3 = lower(t2)
    t4 = remove_punct(t3)
    t5 = remove_stopwords(t4)
    t6 = lemmatize_words(t5)
    return t6
```

```
In [4]: tweet_df = pd.DataFrame(columns=['tweet', 'sentiment', 'NA'])
df_test = pd.DataFrame(columns=['tweet', 'sentiment', 'NA'])

for file in glob.glob("*.tsv"):
    if 'final_test' in file:
        df_test_cur = read_tsv(file)
        df_test = pd.concat([df_test, df_test_cur])
    else:
        df_train_cur = read_tsv(file)
        tweet_df = pd.concat([tweet_df, df_train_cur])
```

We will print some of the training and test datasets tweets to get an idea of how they look before pre-processing.

```
In [5]: print(tweet_df[['tweet', 'sentiment']])
```

	tweet	sentiment
0	05 Beat it - Michael Jackson - Thriller (25th ...	neutral
1	Jay Z joins Instagram with nostalgic tribute t...	positive
2	Michael Jackson: Bad 25th Anniversary Edition ...	neutral
3	I liked a @YouTube video http://t.co/AaR3pjp2P...	positive
4	18th anniv of Princess Diana's death. I still ...	positive
...
1137	Maybe it was - his - fantasy ?	positive
1138	It was ok , but they always just seem so nervo...	negative
1139	It is streamable from YepRoc -- matter of fact...	positive
1140	comment telling me who you are , or how you fo...	positive
1141	im on myspace ... ill try and find you and add...	neutral

[53368 rows x 2 columns]

```
In [6]: print(df_test[['tweet', 'sentiment']])
```

```

                                tweet sentiment
0      #ArianaGrande Ari By Ariana Grande 80% Full ht...  neutral
1      Ariana Grande KIIS FM Yours Truly CD listening...  positive
2      Ariana Grande White House Easter Egg Roll in W...  positive
3      #CD #Musics Ariana Grande Sweet Like Candy 3.4...  positive
4      SIDE TO SIDE 🥰 @arianagrande #sidetoside #aria...  neutral
...
11901  @dansen17 update: Zac Efron kissing a puppy ht...  positive
11902  #zac efron sex pic skins michelle sex https://...  neutral
11903  First Look at Neighbors 2 with Zac Efron Shirt...  neutral
11904  zac efron poses nude #lovely libra porn https:...  neutral
11905  #Fashion #Style The Paperboy (NEW Blu-ray Disc...  neutral

[11906 rows x 2 columns]
```

We will then pre-process the tweets and select a subset of the tweets so that they work with our specified number of batches.

```
In [7]: import string
        from nltk.corpus import stopwords

        clean_tweets = []
        for tweet in tweet_df['tweet']:
            clean_tweets.append(clean_tweet(tweet))

        clean_tweets_1 = clean_tweets[:53000]

        clean_tweets_test = []
        for tweet in df_test['tweet']:
            clean_tweets_test.append(clean_tweet(tweet))

        clean_tweets_test_1 = clean_tweets_test[:11900]

        print("original length of train tweets: ", len(clean_tweets))
        print("original length of test tweets: ", len(clean_tweets_test))

        print("custom length of train tweets: ", len(clean_tweets_1))
        print("custom length of test tweets: ", len(clean_tweets_test_1))
```

```
original length of train tweets:  53368
original length of test tweets:   11906
custom length of train tweets:   53000
custom length of test tweets:    11900
```

```
In [8]: print(clean_tweets_1[:10])
print(clean_tweets_test_1[:10])
```

```
[' beat michael jackson thriller th anniversary edition hd', 'jay
z joins instagram nostalgic tribute michael jackson jay z apparent
ly joined instagram saturday ', 'michael jackson bad th anniversar
y edition picture vinyl unique picture disc vinyl includes origina
l ', ' liked youtube video one direction singing man mirror mich
ael jackson atlanta ga june ', 'th anniv princess dianas death st
ill want believe living private island away public michael j
ackson', 'oridaganjazz st time heard michael jackson sing honolu
lu hawaii restaurant radio abc loved ', 'michael jackson ap
peared saturday th place top miamis trends trndnl', ' old en
ough remember michael jackson attending grammys brooke shields w
ebster sat lap show', 'etbrowser u enjoy nd rate michael jackso
n bit honest ques like cant feel face song god obvious want mj
', ' weeknd closest thing may get michael jackson long timeesp
ecially since damn near mimics everything']
['arianagrande ari ariana grande full singer actress', 'ariana gra
nde kiis fm truly cd listening party burbank arianagrande', 'arian
a grande white house easter egg roll washington arianagrande', 'cd
musics ariana grande sweet like candy oz ml sealed box authentic
new', 'side side 🥺 arianagrande sidetoside arianagrande musically
comunidadgay lgbt🌈 lotb...', 'hairspray live previews macys thanksg
iving day parade arianagrande televisionnbc', 'lindsaylohan 'feelin
g thankful' blasting arianagrande wearing 'toomuch...', ' hate lo
ve songs dammit arianagrande', 'ariana grande [right ft big sean]
アリアナ arianagrande', ' one would prefer listen whole day 🥰👉
could never choose arianagrande intoyou sidetoside songs poll']
```

Then, we will convert our data to integer sequences (where each tweet is represented as a sequence of numbers and the numbers are the index of the words in the whole vocabulary). This prepared our data to be fed into the embedding layer.

```
In [9]: # converting tweets to integer sequences
tokenizer = Tokenizer()
tokenizer.fit_on_texts(clean_tweets_1)
tweet_sequences = tokenizer.texts_to_sequences(clean_tweets_1)
word_index = tokenizer.word_index
print('Found %s unique words in train tweets.' % len(word_index))

# converting tweets to integer sequences
tweet_sequences_test = tokenizer.texts_to_sequences(clean_tweets_test_
1)
```

Found 66956 unique words in train tweets.

We split the training data into train and validation so we have a better idea of how our model is doing on the validation set in each epoch.

```
In [10]: #preparing train lables
tweet_df.loc[tweet_df.sentiment == "positive", "sentiment"] = 2
tweet_df.loc[tweet_df.sentiment == "neutral", "sentiment"] = 1
tweet_df.loc[tweet_df.sentiment == "negative", "sentiment"] = 0

labels = tweet_df["sentiment"].tolist()
labels = [ int(x) for x in labels ]

labels_1 = labels[:53000]

# Split the train set into train and validation set
x_train, x_val, y_train, y_val = train_test_split(tweet_sequences, labels_1, test_size = 0.2, random_state=17)

#preparing test labels
df_test.loc[df_test.sentiment == "positive", "sentiment"] = 2
df_test.loc[df_test.sentiment == "neutral", "sentiment"] = 1
df_test.loc[df_test.sentiment == "negative", "sentiment"] = 0

labels_test = df_test["sentiment"].tolist()
labels_test = [ int(x) for x in labels_test ]

labels_test_1 = labels_test[:11900]

print("length of x train ", len(x_train))
print("length of y train ", len(y_train))
print("length of x val: ", len(x_val))
print("length of y val: ", len(y_val))

length of x train  42400
length of y train  42400
length of x val:   10600
length of y val:   10600
```

To further prepare our data to be fed into the embedding layer, we convert the list of integer sequences to tensors and we will pad 0s and the end of each sequence so that our samples are all of equal size.


```
In [11]: # Training batch size
batch_size = 100

# Put into tensors
x_train = [torch.tensor(x) for x in x_train]
X_train = nn.utils.rnn.pad_sequence(x_train, batch_first=True, padding_value=0).long()
X_train = X_train.view(-1, batch_size, X_train.shape[1])

x_val = [torch.tensor(x) for x in x_val]
X_val = nn.utils.rnn.pad_sequence(x_val, batch_first=True, padding_value=0).long()
X_val = X_val.view(-1, batch_size, X_val.shape[1])

x_test = [torch.tensor(x) for x in tweet_sequences_test]
X_test = nn.utils.rnn.pad_sequence(x_test, batch_first=True, padding_value=0).long()
X_test = X_test.view(-1, batch_size, X_test.shape[1])

y_train = torch.tensor(y_train).view(-1, batch_size)
y_val = torch.tensor(y_val)
y_test = torch.tensor(labels_test_1)

# Apply the same length of X_train on X_val and X_test
len_voc = int((X_train.max()+1).item())
```

```
In [12]: # double checking the shapes

print("shape of X train ", X_train.shape)
print("shape of y train ", y_train.shape)

print("shape of X val ", X_val.shape)
print("shape of y val ", y_val.shape)

print("shape of X test ", X_test.shape)
print("shape of y test ", y_test.shape)

shape of X train  torch.Size([424, 100, 555])
shape of y train  torch.Size([424, 100])
shape of X val    torch.Size([106, 100, 626])
shape of y val    torch.Size([10600])
shape of X test   torch.Size([119, 100, 653])
shape of y test   torch.Size([11900])
```

```
In [13]: # pytorch model layout inspired from https://towardsdatascience.com/convolutional-neural-network-in-natural-language-processing-96d67f91275c

import torch.optim as optim

# The shape of our embeddings will be the size of words in the vocab,
# followed by the dimension of embedding (50)
embeddings = nn.Embedding(len_voc, 50)

# Build CNN model
class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.embeddings = nn.Embedding(len_voc, 50)
        #creating a convolution with 1 inout channel, 100 output channels, and 3x 50 kernel size
        self.cnn = nn.Conv2d(1, 100, (3, 50))

        # we will have three neurons at the last layer before the activation function for our three classes.
        self.clf = nn.Linear(100, 3)

    def forward(self, x):
        x = self.embeddings(x)
        # Add an extra dimension for CNN
        x = x.unsqueeze(1)
        # Apply CNN
        x = self.cnn(x)
        # Choose the maximum value of each filter and delete the extra dimension
        x = x.max(2)[0].squeeze(2)
        # Choose the most important features for the classification
        x = F.relu(x)
        # Apply linear nn for classification
        x = self.clf(x)
        return x

model = Model()
optimizer = optim.Adam(model.parameters(), lr=1e-3)
criterion = nn.CrossEntropyLoss()
```

```
In [14]: from sklearn import metrics
from collections import Counter

# Function for evaluating returns f1 followed by accuracy
def get_f1(X, y_real):
    y_pred = []
    for x in X:
        # Choose the value (class label) with higher probability
        predicted = model(x).argmax(1).cpu().detach()
        y_pred.append(predicted)

    y_pred_torch = torch.cat(y_pred)
    acc = metrics.accuracy_score(y_true=y_real, y_pred= y_pred_torch)
    return metrics.f1_score(y_true=y_real, y_pred=y_pred_torch, average=
'micro'), acc, y_pred
```

```
In [16]: # Training steps
epochs = 20
LOSS = []
for e in range(epochs):
    for i, (x, y) in enumerate(zip(X_train, y_train)):

        # Delete the previous values of the gradient
        optimizer.zero_grad()
        x, y = x, y

        y_pred = model(x)
        loss = criterion(y_pred, y)

        # Compute the gradient
        loss.backward()

        # Apply the optimization method for one step
        optimizer.step()

        LOSS.append(loss.item())
        if i%200==0:
            with torch.no_grad():
                f1, acc, y_pred = get_f1(X_val, y_val)
                print('Epoch: %d \t Batch: %d \t Loss: %.10f \t F1_val: %.
10f \t Accuracy: %.10f'%(e,i, torch.tensor(LOSS[-100:]).mean(), f1, ac
c))
```

Epoch: 0	Batch: 0	Loss: 0.9495213032	F1_val: 0.4
471698113	Accuracy: 0.4471698113		
Epoch: 0	Batch: 200	Loss: 0.9902819991	F1_val: 0.5
126415094	Accuracy: 0.5126415094		
Epoch: 0	Batch: 400	Loss: 0.9411697388	F1_val: 0.5
493396226	Accuracy: 0.5493396226		
Epoch: 1	Batch: 0	Loss: 0.9342650771	F1_val: 0.5
516981132	Accuracy: 0.5516981132		
Epoch: 1	Batch: 200	Loss: 0.8722500801	F1_val: 0.5
641509434	Accuracy: 0.5641509434		
Epoch: 1	Batch: 400	Loss: 0.8265760541	F1_val: 0.5
750000000	Accuracy: 0.5750000000		
Epoch: 2	Batch: 0	Loss: 0.8226707578	F1_val: 0.5
768867925	Accuracy: 0.5768867925		
Epoch: 2	Batch: 200	Loss: 0.7616884112	F1_val: 0.5
855660377	Accuracy: 0.5855660377		
Epoch: 2	Batch: 400	Loss: 0.7139935493	F1_val: 0.5
846226415	Accuracy: 0.5846226415		
Epoch: 3	Batch: 0	Loss: 0.7108892798	F1_val: 0.5
871698113	Accuracy: 0.5871698113		
Epoch: 3	Batch: 200	Loss: 0.6455029249	F1_val: 0.5
943396226	Accuracy: 0.5943396226		
Epoch: 3	Batch: 400	Loss: 0.6001831293	F1_val: 0.5
927358491	Accuracy: 0.5927358491		
Epoch: 4	Batch: 0	Loss: 0.5976411700	F1_val: 0.5
944339623	Accuracy: 0.5944339623		
Epoch: 4	Batch: 200	Loss: 0.5271141529	F1_val: 0.5
950943396	Accuracy: 0.5950943396		

```

-----
KeyboardInterrupt                                Traceback (most recent call
last)
<ipython-input-16-73c4c72a6f97> in <module>
      9         x, y = x, y
     10
--> 11         y_pred = model(x)
     12         loss = criterio(y_pred, y)
     13

~/opt/anaconda3/envs/neural_networks/lib/python3.7/site-packages/tor
ch/nn/modules/module.py in _call_impl(self, *input, **kwargs)
     725         result = self._slow_forward(*input, **kwargs)
     726     else:
--> 727         result = self.forward(*input, **kwargs)
     728     for hook in itertools.chain(
     729         _global_forward_hooks.values(),

<ipython-input-13-fd4d2e4c7af0> in forward(self, x)
     22         x = x.unsqueeze(1)

```

```

23         # Apply CNN
--> 24         x = self.cnn(x)
25         # Choose the maximum value of each filter and delete
the extra dimension
26         x = x.max(2)[0].squeeze(2)

~/opt/anaconda3/envs/neural_networks/lib/python3.7/site-packages/tor
ch/nn/modules/module.py in _call_impl(self, *input, **kwargs)
725         result = self._slow_forward(*input, **kwargs)
726     else:
--> 727         result = self.forward(*input, **kwargs)
728     for hook in itertools.chain(
729         _global_forward_hooks.values(),

~/opt/anaconda3/envs/neural_networks/lib/python3.7/site-packages/tor
ch/nn/modules/conv.py in forward(self, input)
421
422     def forward(self, input: Tensor) -> Tensor:
--> 423         return self._conv_forward(input, self.weight)
424
425 class Conv3d(_ConvNd):

~/opt/anaconda3/envs/neural_networks/lib/python3.7/site-packages/tor
ch/nn/modules/conv.py in _conv_forward(self, input, weight)
418         _pair(0), self.dilation, self.gr
oups)
419         return F.conv2d(input, weight, self.bias, self.strid
e,
--> 420             self.padding, self.dilation, self.gr
oups)
421
422     def forward(self, input: Tensor) -> Tensor:

KeyboardInterrupt:

```

```

In [17]: f1, acc, y_pred = get_f1(X_test, y_test)
print(" f1 is ", f1, " accuracy is ", acc)

```

```

f1 is  0.5085714285714286  accuracy is  0.5085714285714286

```

```
In [18]: from sklearn.metrics import classification_report
predicted_labels_list = []

flat_list = [item for sub in y_pred for item in sub]

for elem in flat_list:
    predicted_labels_list.append(elem.tolist())

print(len(predicted_labels_list))
print(classification_report(labels_test[:11900], predicted_labels_list
))
```

```
11900
```

	precision	recall	f1-score	support
0	0.55	0.28	0.37	3811
1	0.55	0.66	0.60	5739
2	0.39	0.52	0.45	2350
accuracy			0.51	11900
macro avg	0.50	0.48	0.47	11900
weighted avg	0.52	0.51	0.49	11900

In []: